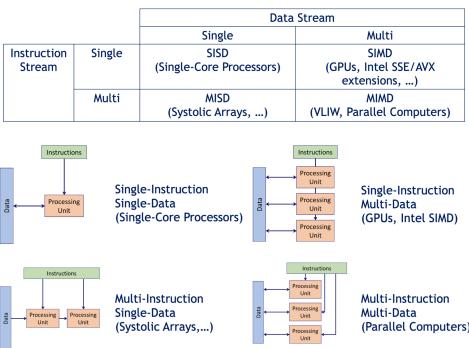


## Instruction Type (SIMD = vector dot prod)



## Intel AVX 512 registers

XMM0-15 (128), YMM0-15 (256), ZMM0-31 (256)  
Does 32x 8 bits operations in one instruction

### Instruction types

- **Pointwise:** saxpy (ax+y), add, scale
  - **Reduction:** sum, average, min, max
- ```
#pragma omp parallel for reduction([op]:[var])
```

### Performance bound types

Which takes longer, I/O or the algorithm itself

- **Memory:** larger complexity for I/O
- **Compute:** larger complexity for the algorithm itself

### Laplacian matrix

Treats convolution as a (sparse) matrix multiplication

### Conjugate gradient operation

Solves for  $\mathbf{Ax} = \mathbf{b}$  given a solution exists, allows systems of linear equation to be solved at once ( $\mathbf{r} = \mathbf{b} - \mathbf{Ax}$ )

- **Pros:** more compact and cost effective for dense  $\mathbf{A}$
- **Cons:** space inefficient if  $\mathbf{A}$  is sparse

## CSR (Compressed Sparse Row) Matrix

$$\mathbf{A} = \begin{pmatrix} 10 & 11 & 12 \\ & 13 & 14 \\ 15 & & 16 & 17 \\ & 18 & & 19 \end{pmatrix}$$

```
int offsets[] = { 0, 3, 5, 8, 10 }
int col[] = { 0, 1, 3, 1, 2, 0, 2, 3, 1, 3 }
int value[] = { 10, 11, 12, 13, 14, 15, 16, 17, 18, 19 }
len(offsets) = n+1 starts with 0 and ends with len(value).
```

- **Pros:** very space efficient for sparse matrices

- **Cons:** does nothing if  $\mathbf{A}$  is full/dense

```
for (int i = 0; i < N; i++)
    y[i] = 0;
    for (int k = rowOffsets[i]; k < rowOffsets[i+1]; k++)
        const int j = columnIndices[k];
        y[i] += values[k] * x[j];
```

mkl\_cspblas\_scsrgemv multiplies CSR matrix with vector

### GEMM operation

Does matrix-matrix multiplication ( $C = AB$ ) as blocks

Aligns the entries since matrix is **ROW MAJOR**

Transposes  $B$  to make it **COLUMN MAJOR**

```
for (int ii = 0; ii < BLOCK_SIZE; ii++)
    for (int jj = 0; jj < BLOCK_SIZE; jj++)
        #pragma omp simd
            for (int kk = 0; kk < BLOCK_SIZE; kk++)
                localC[ii][jj] += localA[ii][kk] * localB[jj][kk];
```

Complexity:  $O(N^3)$  but better spatial locality

**Small block:** better locality, wastes space, data bound

**Large block:** doesn't fit in cache, compute bound

Use  $i, k, j$  instead of  $i, j, k$  if  $B$  is not transposed

```
for (int ii = 0; ii < BLOCK_SIZE; ii++)
    for (int kk = 0; kk < BLOCK_SIZE; kk++) {
        #pragma omp simd
            for (int jj = 0; jj < BLOCK_SIZE; jj++)
                bC[ii][jj] += bA[ii][kk] * bB[kk][jj];
```

## Sparse Operations

I.e. saxpy  $y[offset[i]] += a*x[offset[i]]$

**offset** can be stored in blocks **blockOffset**

```
#pragma omp parallel for
    for (int b = 0; b < blockOffsets.size(); b++) {
        _mm_prefetch (&x[b*blockOffsets[b]+L_PREFETCH_DISTANCE], _MM_HINT_T2 );
        _mm_prefetch (&x[b*blockOffsets[b]+L_PREFETCH_DISTANCE], _MM_HINT_T1 );
        _mm_prefetch (&y[b*blockOffsets[b]+L_PREFETCH_DISTANCE], _MM_HINT_T2 );
        _mm_prefetch (&y[b*blockOffsets[b]+L_PREFETCH_DISTANCE], _MM_HINT_T1 );
    #pragma omp end parallel
        for (int i = 0; i < BLOCK_SIZE; i++)
            y[blockOffsets[b]+i] += scale * x[blockOffsets[b]+i];
```

### Struct of array stores x,y,z separately

- **Pros:** allows SIMD, better locality per channel, easier to resize one of the channels

- **Cons:** less compatibility with SIMD for sparse, uses multiple localities (one per channel), concerns with vector type (i.e. coordinates), TLB pressure

**Array of struct** arr[0]=(x[0],y[0],z[0]) allows saxpy to be done on offset directly

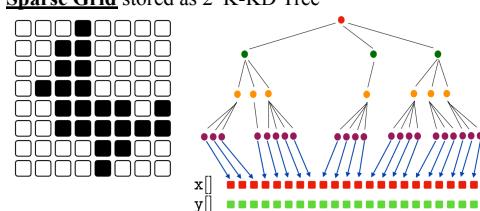
- **Pros:** more natural access, better locality if all channels are accessed, exploits SIMD if sparse

- **Cons:** waste access cost if one channel is accessed, alignment issue if not  $2^n$ , SIMD problem if not sparse

### Arrays of structs of arrays (hybrid of the two)

```
data[] ━━━━ ━━━━ ━━━━ ━━━━ ━━━━ ━━━━ ━━━━ ━━━━
```

### Sparse Grid stored as $2^K$ -KD Tree



**Problem:** neighboring data can be from different tree

## OpenVDB/NanoVDB

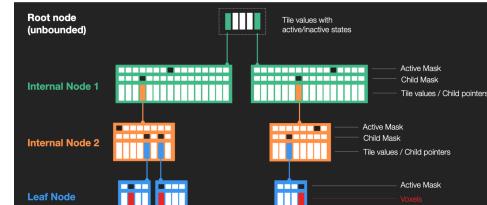
Stores 8x8x8 block for a leaf node instead of one point

Uses caching structure with easy neighbor access

**IndexGrid:** tells VDB what the coordinate range is.

GridBuilder creates a tree. IndexGrid also stored data

- **Pros:** works better for dense matrices since nearby values are stored in the same block. Also skips sparse parts if FLAG is set to INACTIVE

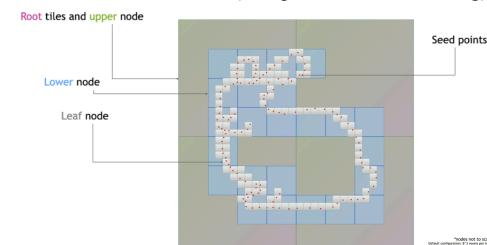


- **Root:** hash table or linear array (store as needed)

- **Upper (1):** store 32x32x32 lower nodes

- **Lower (2):** store 16x16x16 leaf nodes

- **Leaf node:** stores 512 x (data point +1 bit active flag)



OpenVDB stores tree as pointers, NanoVDB stores the tree in a flattened array (+offset)

firstPrivate([var]) var is created in all threads but only the first thread's value is copied to master. Copies the first because VDB can alter data entries

## Assembly Intrinsics (Built-in ASM)

```
#include "immintrin.h"

for (int ii = 0; ii < BLOCK_SIZE; ii++)
    for (int kk = 0; kk < BLOCK_SIZE; kk++)
        for (int jj = 0; jj < BLOCK_SIZE; jj += nW) {
            __m256 vB = _mm256_load_ps(&B[B*kk][jj]);
            __m256 vA = _mm256_set1_ps(bA[ii][kk]);
            __m256 vC = _mm256_fmadd_ps(vA, vB, vC);
            _mm256_store_ps(&bC[ii][jj], vC);
        }
```

nW=16 number of floats to go into SIMD kernel call

### Solving Linear Equations with Triangular L

$Lx=b$  easy to solve if  $L$  lower triangular ( $\text{cblas}_?trsm$ )

**Forward substitution** method, solves  $x_1, x_2, \dots, x_m$

$$\begin{aligned} \ell_{1,1}x_1 &= b_1 & x_1 &= \frac{b_1}{\ell_{1,1}}, \\ \ell_{2,1}x_1 + \ell_{2,2}x_2 &= b_2 & x_2 &= \frac{b_2 - \ell_{2,1}x_1}{\ell_{2,2}}, \\ &\vdots & \vdots & \vdots \\ \ell_{m,1}x_1 + \ell_{m,2}x_2 + \dots + \ell_{m,m}x_m &= b_m & x_m &= \frac{b_m - \sum_{i=1}^{m-1} \ell_{m,i}x_i}{\ell_{m,m}}. \end{aligned}$$

**BLAS level 3** direct solver for  $AX=B$  if  $A$  is dense and symmetric (saves bandwidth),  $X$  and  $B$  are matrices. GEMM is a part of it.

**LAPACK** solve  $AX = B$ , preprocesses  $A$  into more manageable forms (prefer  $A$  is lower triangular matrix or symmetric matrix)

- **Iterative:** easier to set up and no preprocess required, but doesn't guarantee convergence and is less accurate
- **Direct:** preprocesses  $A$ , easy to parallel if  $A$  is dense, always work, but more expensive up to  $O(N^3)$

**PARDISO:** direct solver if  $A$  is sparse. CSL/CSR format

- 1 reorder  $A$  into a favorable form (more diagonal)
- 2 Cholesky decomposition  $A = LL^*$ , faster if  $L$  is more sparse, is very serial (like Gaussian elimination)

Red rect in  $A$  also exists in  $L$

- 3 solve the equation since  $L$  is now triangular

## Convolution

```
for (int i = 0; i < Nz; i += 8)
    for (int j = 0; j < Nz; j += 8)
        for (int k = 0; k < Nz; k += 8) {
```

```
    float inLocal[10][10][10][Di];
    for (int di = -1; di < 8; di++)
        for (int dj = -1; dj < 8; dj++)
            for (int dk = -1; dk < 8; dk++)
                for (int inDim = 0; inDim < Di; inDim++)
                    inLocal[di+1][dj+1][dk+1][inDim] = in[i+di][j+dj][k+dk][inDim];

    using in_local_tensor_array_t = float (&)[10][10][10][Di];
    inLocal_tensor_array_t inLocalShifted = reinterpret_cast<in_local_tensor_array_t>(inLocal[1][1][1][0]);
    float outLocal[8][8][8][Do];
    outLocal[0][0][0][0] = 0;

    localStencilOpDi.Do(inLocalShifted, outLocal, stencil);
```

```
    for (int di = 0; di < 8; di++)
        for (int dj = 0; dj < 8; dj++)
            for (int dk = 0; dk < 8; dk)
                for (int outDim = 0; outDim < Do; outDim++)
                    out[i+di][j+dj][k+dk][outDim] += outLocal[di][dj][dk][outDim];
```

Optimized version, each 8x8x8 region of output is computed from a 10x10x10 region of input. localStencil does convolution naively

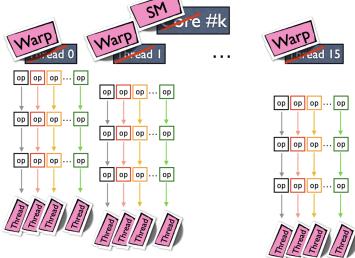
```
using in_tensor_array_t = float (&)[Nx][Ny][Nz][Di];
using out_tensor_array_t = float (&)[Nx][Ny][Nz][Do];
using stencil_array_t = float (&)[3][3][3][Di][Do];
```

```
for (int di = -1; di < 1; di++)
    for (int dj = -1; dj < 1; dj++)
        for (int dk = -1; dk < 1; dk++)
            for (int i = 0; i < 8; i++)
                for (int j = 0; j < 8; j++)
                    for (int k = 0; k < 8; k++)
                        for (int inDim = 0; inDim < Di; inDim++)
                            for (int outDim = 0; outDim < Do; outDim++)
                                out[i][j][k][outDim] +=
                                    stencil[di][dj][dk][inDim][outDim] *
                                    in[i+di][j+dj][k+dk][inDim];
```

di,dj,dk goes in the outer loop so stencil stays in cache as long as possible (stencil[di][dj][dk] is const)  
Inner 3 loops are effectively GEMM operations on  $in[i+di][j+dj] * stencil[di][dj] \rightarrow out[i][j]$  (becomes a saxpy equivalent)

## Hardware Organization: CPU vs GPU

One core can consist of several threads, per core, one thread runs at a time (scheduling) and share L1-3 caches



### GPU:

- Core → **Streaming Multiprocessor (SM)**, tens of SM per GPU with its own cache
- Independent stream → **warp**, each capable of doing 32-wide SIMD instructions
- 1 SM can have up to 32 Warps, warp runs in round robin or scheduler (literally like CPU schedulers)
- $32 \times 32 = 1024$  streams per SM, these are called **threads**
- threads from the same warp are dependent if SIMD is used (cannot execute different instructions)
- using **SIMT**, each thread can run different instructions

```
__global__
void func() {
    int threadID = threadIdx.x; // index [0..1023] of this thread
    int warpID = threadIdx.x >> 5; // faster way to do threadID % 32
    if (warpID == 0)
        doFirstWarpStuff();
    else
        doOtherWarpStuff();
}
```

- **thread block** = divide all threads into blockid/threadid

```
__global__
void myKernel(float *array) { // Kernel on GPU
    int threadID = threadIdx.x;
    int blockID = blockIdx.x;
    int elementID = blockID * blockDim.x + threadID;

    array[elementID] += 1.0;
}

#define SIZE 1024*1024

void myCPUFunc(){
    float *myArray = ...; // Already in GPU memory
    int blockDim = 512; // thread block size
    int gridDim = SIZE/blockDim;
    myKernel<<<gridDim,blockDim>>>(&myArray[0]);
}
```

- **higher dimension case:**

```
__device__
float GPUfunc(float x) { return x*x*cos(x); }

__global__
void kernel(float (&u)[Nx][Ny], float (&Lu)[Nx][Ny]) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;
    Lu[i][j] = GPUfunc(u[i][j]);
}

void myCPUFunc(){
    float myU[Nx][Ny] = ...; // Already in GPU memory
    float myLU[Nx][Ny] = ...; // Already in GPU memory
    dim3 blockDim(4,4,1); // use "1" for unused dims
    dim3 gridDim(Nx/4,Ny/4,1);
    Kernel<<<gridDim,blockDim>>>(myU, myLU);
}
```

- **\_\_global\_\_** can be called from anywhere

- **\_\_device\_\_** can be called only from GPU

- **\_\_host\_\_** can be called only from CPU

- only arguments are copied from CPU → GPU

`cudaMemcpyHostToDevice: Copy CPU->GPU`

`cudaMemcpyDeviceToHost: Copy GPU->CPU`

`cudaMemcpyDeviceToDevice: Copy GPU->GPU`

- `cudaMalloc` and `cudaMemcpy` allocates and copies

đến operation **memmove** gán **nhà** **nhà**

卧槽

```
// Device code
__global__ void VecAdd(float* A, float* B, float* C, int N)
{
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    C[i] = A[i] + B[i];
}

// Host code
int main()
{
    int N = ...;
    size_t size = N * sizeof(float);

    // Allocate input vectors h_A and h_B in host memory
    float* h_A = (float*)malloc(size);
    float* h_B = (float*)malloc(size);
    float* h_C = (float*)malloc(size);

    // Initialize input vectors
    ...

    // Allocate vectors in device memory
    float* d_A;
    cudaMalloc(&d_A, size);
    float* d_B;
    cudaMalloc(&d_B, size);
    float* d_C;
    cudaMalloc(&d_C, size);

    // Copy vectors from host memory to device memory
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);

    // Invoke kernel
    int threadsPerBlock = 256;
    int blocksPerGrid = (N + threadsPerBlock - 1) / threadsPerBlock;
    VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A, d_B, d_C, N);

    // Copy result from device memory to host memory
    // h_C contains the result in host memory
    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

    // Free device memory
    cudaFree(d_A); cudaFree(d_B); cudaFree(d_C);
}
```

### Inter Thread Communication

Shared memory: user managed cache that allows threads other, there is 64KB of shared memory per SM. Shared between blocks when device code is called

```
#define ROWS 65536
#define COLS 64
// Device code
__global__ void reverseRowsKernel(int (&data)[ROWS][COLS])
{
    int threadID = threadIdx.x;
    int rowID = blockIdx.x;
    __shared__ tempData[COLS];
    auto &rowData = data[rowID];
    // type of rowData is int (&)[COLS]
    tempData[threadID] = rowData[threadID];
    // copy to shared memory
    __syncthreads();
    // make sure all threads have written their part
    rowData[threadID] = tempData[COLS-1-threadID]; // reverse
}
```

- in the code above, `__syncthreads()` awaits all threads to read to `tempData` (on the same BLOCK)

```
__global__ void sumRowsKernel(int (&data)[COLS], int (&sums)[ROWS])
{
    int threadID = threadIdx.x;
    int warpID = threadID >> 5;
    int threadInWarpID = threadID & 0x1f; // Number of thread in warp
    constexpr int nWarps = COLS/32; // Warps in block, here nWarps=2

    __shared__ partialSums[nWarps]; // put only partial sums in smem
    auto &rowData = data[rowID]; // type of rowData is int (&)[COLS]
    auto &sum = sums[rowID];

    int partialSum = __reduce_add_sync(0xffffffff, rowData[threadID]);
    partialSums[warpID] = partialSum; // all threads in warp copy the same value to shared mem (this is ok)

    __syncthreads(); // make sure all threads have written their part

    if (threadID == 0) {
        sum = 0;
        for (int warp = 0; warp < nWarps; warp++)
            sum += partialSums[warp]; // much less work to add up
    }
}
```

- `__reduce_add_sync` does reduction on the same WARP

```
__global__ void sumRowsKernel(int (&data)[COLS], int &totalSum)
{
    int threadID = threadIdx.x;
    int rowID = blockIdx.x;

    __shared__ tempData[COLS];
    auto &rowData = data[rowID]; // type of rowData is int (&)[COLS]
    tempData[threadID] = rowData[threadID]; // copy to shared memory

    __syncthreads(); // make sure all threads have written their part

    if (threadID == 0) {
        int blocksum = 0;
        for (int col = 0; col < COLS; col++)
            blocksum += tempData[col];
        atomicAdd(&totalSum, blocksum);
    }
    __syncthreads();

    // Host code
    void myCPUcode()
    {
        int data[ROWS][COLS] = ... // Presume this is already on GPU
        int totalSum; // Presume this is on GPU, and initialized to zero
        sumRowsKernel<<<ROWS,COLS>>>(data,totalSum);
    }
}
```

- `atomicAdd` adds `blocksum` (local) to `totalSum` (shared)



「がんばろうね！」ミクちゃんと言つていまし  
た



พีເລຍ໌ພົມຄວຍດຸຍ໌ ພິເປ່າ  
ຂລຸ່ມກະຄວຍເປົາແຕ



早上好中国 现在我  
有冰激淋